

Comparing Quantization Across Hardware

Benchmarking LLM Inference on Edge Devices

Mike Doan · Jenny Dinh

CSC 4228 & 6228 – Security in IoT

Georgia State University

Fall 2024

Doan & Dinh | CSC 4228/6228 | Fall 2024

Introduction & Background

The Rise of Local LLM Deployment

- Growing interest in running LLMs **locally** rather than via cloud APIs
- Target environments:
 - **Edge devices**: smartphones, embedded systems
 - **Consumer hardware**: laptops, gaming PCs
 - **Enterprise GPUs**: data center deployments

The Memory Challenge

- LLMs contain **massive weights** requiring significant memory

- Activations during inference are also substantial
- Example: Llama-3.1-8B requires 16GB in FP16

Background: What is Quantization?

Quantization reduces the precision of model weights and activations to decrease memory footprint and improve inference speed.

Precision	Bits per Weight	Memory (8B model)
FP16 / BF16	16 bits	16 GB
INT8 (W8A8)	8 bits	8 GB
INT4 (W4A8)	4 bits	4 GB
NF4	4 bits	4 GB

Trade-off: Lower precision → smaller memory, faster inference, but potential accuracy loss

Background: Quantization Strategies

Weight-Only Quantization

- INT8: 8-bit integer weights
- INT4: 4-bit integer weights
- NF4: 4-bit NormalFloat (QLoRA)

KV-Cache Quantization

- Reduces memory for long contexts
- INT8 or INT4 KV cache

Activation Quantization

- W8A8: Both weights and activations in INT8
- W4A4: Aggressive 4-bit for both

Group-wise Quantization

- **GPTQ**: Post-training quantization
- **AWQ**: Activation-aware weights

Purpose & Motivation

Research Questions

1. How do different **quantization strategies** impact inference performance across hardware?
2. What are the **accuracy-efficiency trade-offs** for edge deployment?
3. Can we create a **reproducible benchmarking framework** for the community?

Gap in Existing Work

- Few plug-and-play toolkits for **cross-platform quantization sweeps**
- Limited public data on **edge device performance**
- Need for standardized comparison methodology

Target Hardware & Scope

Edge Devices

- Google Pixel 7/8 (Android)
 - ARM CPU + Adreno GPU
 - OpenCL backend

Runtimes Evaluated

- `llama.cpp` (GGUF format)
- `tinygrad` (Metal/OpenCL)
- MLC-LLM (TVM-based)

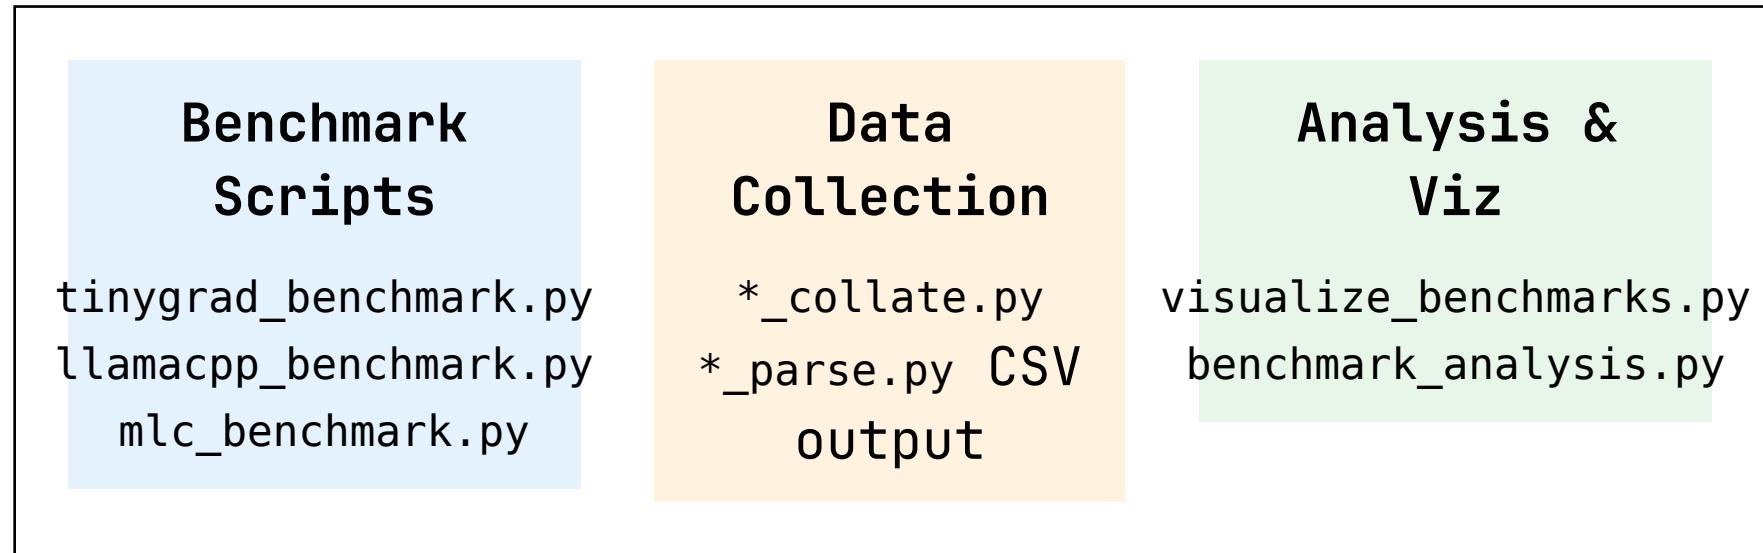
Consumer Hardware

- MacBook (Apple Silicon)
 - Metal backend
 - Unified memory

Models

- Llama-3.2-1B-Instruct
- Qwen2.5-1.5B
- (Extensible to larger models)

System Architecture



Key Design Principles

- **Reproducibility:** Seeded runs, UUID tracking

- **Extensibility:** Easy to add new backends/devices
- **Standardized schema:** Consistent CSV format across all backends

Data Schema

Each benchmark run captures:

BenchmarkRow:

```
step: int                      # Token generation step
enqueue_latency_ms: float      # Time to enqueue operation
total_latency_ms: float        # Total time for step
tokens_per_sec: float          # Throughput
memory_throughput_gb_s: float
param_throughput_gb_s: float
platform: str                  # Darwin, Linux, Android
device: str                     # Metal, OpenCL, CUDA
hostname: str                   # Machine identifier
size: str                       # Model size (1B, 3B, 8B)
quantize: str                   # nf4, int8, float16, default
seed: int                        # For reproducibility
uuid: str                        # Unique run identifier
```

Implementation: Tinygrad Backend

- Pure Python ML framework with Metal/OpenCL support
- Server mode for OpenAI-compatible API

```
# Start inference server
PYTHONPATH=./deps/tinygrad/ python tinygrad_benchmark.py \
--port 7776 --size 1B
```

Quantization Support

- default: No quantization
- nf4: 4-bit NormalFloat
- int8: 8-bit integer
- float16: Half precision

Implementation: llama.cpp Backend

- C/C++ implementation with GGUF model format
- Highly optimized for CPU and GPU inference

```
# Run benchmark  
python llamacpp_benchmark.py
```

```
# Collate results  
python llamacpp_collate.py
```

Key Features

- Native quantization support in model files
- Metal backend for Apple Silicon

- OpenCL for Android devices

LLM Evaluation: Verifiers Framework

Beyond raw throughput, we measure **downstream task accuracy**:

Setup

```
# Start model server
python tinygrad_benchmark.py \
    --port 7776 --size 1B
```

```
# OpenAI proxy (streaming)
python openai_proxy.py \
    --backend-port 7776 \
    --proxy-port 7777
```

Run Evaluation

```
# GSM8K math benchmark
OPENAI_API_KEY=dummy \
uv run vf-eval gsm8k \
    -m local \
    -b http://localhost:7777/v1 \
    -n 20 -r 1 -t 512
```

Available benchmarks: gsm8k, math, gpqa, simpleqa, wordle

Metrics Collected

Performance Metrics

Latency

- Time to first token (TTFT)
- Per-token generation time

Throughput

- Tokens per second (tok/s)
- Memory bandwidth (GB/s)

Resource Metrics


Memory

- Model load size (MB)
- Peak utilization

Accuracy

- Downstream benchmark scores
- (GSM8K, MATH, etc.)

Results: Throughput Comparison




Key findings:

- llama.cpp excels with NF4 (46.4 tok/s) and default quantization (41.3 tok/s)
- tinygrad performs best with INT8 (27.5 tok/s) and float16 (25.4 tok/s)
- Performance varies significantly by quantization method

Results: Speedup Analysis

Speedup Comparison vs Tinygrad Baseline




Observations:

- llama.cpp shows 16.6x speedup over tinygrad for NF4 quantization
- tinygrad leads by 8.0x for float16 workloads
- INT8 and default quantizations show modest llama.cpp advantage (1.2-2.5x)

Results: Performance Summary


Performance Summary by Backend



Backend comparison:

- `llama.cpp`: Average 31.0 tok/s, peak 46.4 tok/s (NF4)
- `tinygrad`: Average 18.1 tok/s, peak 27.5 tok/s (INT8)
- Overall `llama.cpp` shows 1.7x better average performance


Results: Quantization Impact



Trends:

- Different backends favor different quantization strategies
- llama.cpp benefits most from aggressive quantization (NF4)
- tinygrad shows more consistent performance across methods

Results: Latency Analysis





Latency variance:

- Box plots show distribution of per-token generation time
- llama.cpp shows lower variance for NF4 and default

- tinygrad has more consistent latency across quantizations


Results: Memory & Parameter Throughput



Observations:

- Memory bandwidth closely correlates with tokens/sec
- Parameter throughput shows compute efficiency per quantization

Results: Multi-Device Comparison



Cross-device insights:

- MacBook (softmacs) shows stronger performance overall
- Android device (localhost) exhibits different optimization characteristics
- Backend choice impacts relative performance across hardware

Results: Performance Summary Table

Backend	Default	FP16	INT8	BF16
llama.cpp	41.3 tok/s	3.2 tok/s	33.2 tok/s	46.4 tok/s
tinygrad	16.6 tok/s	25.4 tok/s	27.5 tok/s	2.8 tok/s

Note: Results from Llama-3.2-1B-Instruct on MacBook (Metal). Actual performance varies with context length and workload.

Challenges Encountered

Software Compatibility

- Different backends for ARM vs x86
- MLC-LLM build complexity
 - “Opted to skip due to finicky build process”

Engineering Challenges

- Creating durable benchmarking suite
- Handling device-specific configurations
 - Pixel 7 OpenCL requires special env vars

Implementation

Inconsistencies

- Quantization standards differ across backends

Android Setup

- SSH into Pixel devices

- Leads to performance discrepancies
- OpenCL library path configuration

Contributions

Reproducible Framework

- Open-source benchmarking toolkit
- Standardized data format
- Easy to extend to new devices

Empirical Results

- Cross-device performance data
- Quantization trade-off analysis
- Downstream accuracy evaluation

Public Codebase

- MIT licensed
- Well-documented
- GitHub repository

Community Benefit

- Others can benchmark their devices

- Append results to shared data bank

Future Work

- **Enterprise GPUs:** Extend benchmarks to A100, H100
- **More Models:** Phi-3 mini, Qwen2.5-3B, Llama-3.1-8B
- **Energy Profiling:** Watt/token measurements
- **KV-Cache Quantization:** INT8/INT4 cache strategies
- **Longer Contexts:** Prefill tokens 256, 1024, 2048

Potential Extensions

- Automated setup script (`curl | sh`)
- CI/CD backend for continuous benchmarking
- Public data aggregation platform

Conclusion

- **Quantization enables practical edge LLM deployment**
 - 4-bit models run on smartphones with acceptable performance
- **Trade-offs are workload-dependent**
 - Memory-constrained? → INT4/NF4
 - Accuracy-critical? → INT8 or FP16
- **Framework enables reproducible research**
 - Standardized benchmarks across diverse hardware
 - Community can contribute additional results

Code: github.com/spikedoanz/t-eai-project

References

1. **llama.cpp** – Georgi Gerganov et al. High-performance LLM inference in C/C++.
<https://github.com/ggerganov/llama.cpp>
2. **tinygrad** – George Hotz et al. A simple deep learning framework.
<https://github.com/tinygrad/tinygrad>
3. **MLC-LLM** – Machine Learning Compilation for LLMs.
<https://github.com/mlc-ai/mlc-llm>
4. **Verifiers** – Prime Intellect. LLM evaluation framework.
<https://github.com/PrimeIntellect-ai/verifiers>
5. **QLoRA** – Dettmers et al. (2023). Efficient Finetuning of Quantized LLMs.
6. **GPTQ** – Frantar et al. (2022). Accurate Post-Training Quantization.
7. **AWQ** – Lin et al. (2023). Activation-aware Weight Quantization.

Thank You!

Questions?

Mike Doan • Jenny Dinh

github.com/spikedoanz/t-eai-project

Doan & Dinh | CSC 4228/6228 | Fall 2024